

Image Classification with Tensorflow

Implementation and Comparison of Deep CNNs for MNIST

Otmane Sabir x Aadil Anil Kumar

Table of Contents

01

Introduction

Brief intro to image classification
& (deep) learning

02

Classification

Implementing NNs using
tensorflow

03

Conclusion

Comparison of model
performances

Introduction

01



Brief intro to image classification & (deep) learning

Image Classification



Process of classifying an image based on visual content



Began with the use of hand-engineered features



Pushed forward by the advent of deep learning



Supervised vs. Unsupervised

Neural Networks



Aims to mimic the
functionality of the brain



Stems from the study of
Artificial Neural Networks



Feedforward vs. Recurrent



Deep networks add more
layers between the input and
output layers

Universal Approximation & the Significance of Depth

NNs work to approximate a function

More layers reduce the chance of overfitting

More layers *can* be computationally cheaper




Figure extracted from Goodfellow's 2014 paper "Multi-digit Number Recognition from Street View Imagery using Deep Convolutional Neural Networks"

MNIST Dataset


Labelled Handwritten Digits,

0 - 9

60,000 Training

10,000 Testing

Grayscale, 28x28



Classification

02

Implementing NNs using tensorflow

Layers Used

Matrix vector
multiplication



Dense



Introduces
non-linearity

Feature
dimensionality
reducer



Flattening



2d matrix to vector,
for input to NN

Sets random input
units to 0



Dropout



Prevents overfitting

Extracts features
from an image



Convolutional



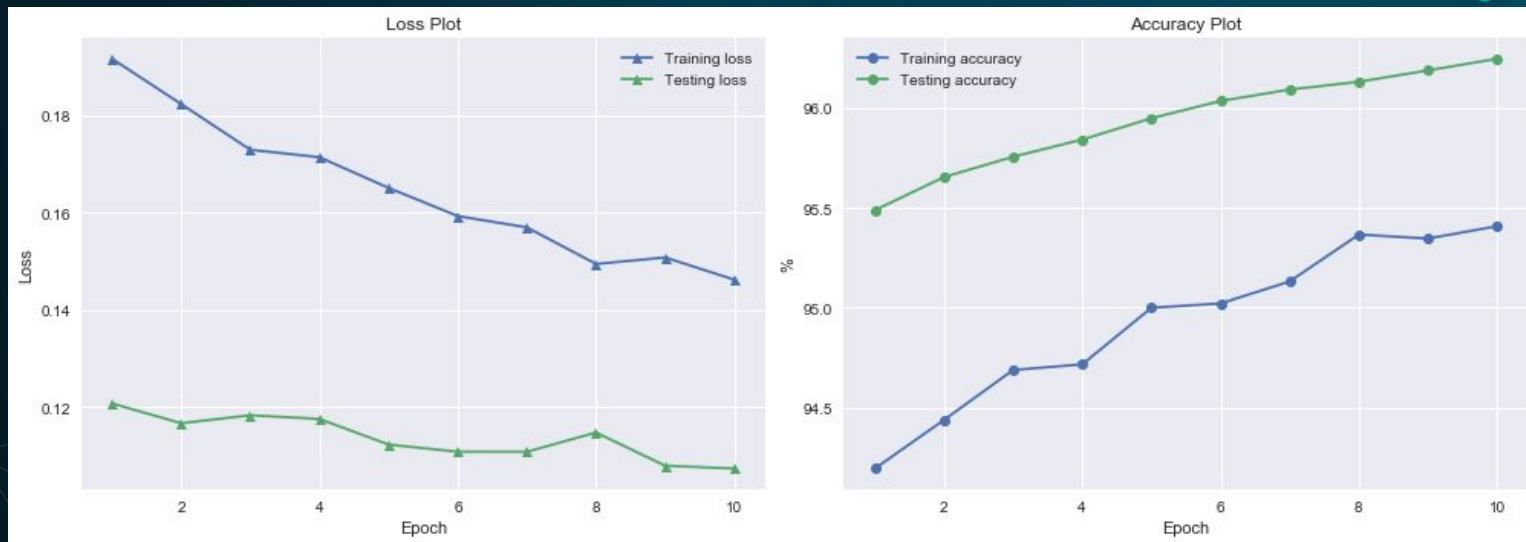
Blurring, sharpening,
edge detection, etc...

- Flattening Layer
- Dense Layer 1 w/ RELU activation
- Dropout Layer 1 w/ 30% dropout rate
- Dense Layer 2 w/ RELU activation



Baseline - Model 1

Results - Model 1




Train Accuracy: 95.4%
Test Accuracy: 96.2%
Trainable Parameters: 36,935

- Flattening Layer
- Dense Layer 1 w/ RELU activation
- Dropout Layer 1 w/ 30% dropout rate
- Dense Layer 2 w/ RELU activation
- Dense Layer 3 w/ RELU activation
- Dense Layer 4 w/ Softmax activation



Adding Layers - Model 2

Results - Model 2



Train Accuracy: 93.2%
Test Accuracy: 94.5%
Trainable Parameters: 38,003




CONVOLUTION

- Convolutional Layer w/ RELU activation
- Flattening Layer
- Dense Layer 1 w/ RELU activation
- Dense Layer 2 w/ Linear activation



Adding Convolutions - Model 3

Results - Model 3




Train Accuracy: 99.9%
Test Accuracy: 98.3%
Trainable Parameters: 2,770,634

Conclusion

03

Comparison of model performances

Model Comparison



Model-1	Model-2	Model-3
Test Accuracy: 96.2%	Test Accuracy: 94.5%	Test Accuracy: 98.3%

References

Herbert, J.. (2019). Machine Learning Lecture Notes. <http://minds.jacobs-university.de/teaching/courses/t2019ml/>

Yi, H., Shiyu, S., Xiusheng, D., & Zhigang, C. (2016). A study on deep neural networks framework. In *2016 IEEE Advanced Information Management, Communicates, Electronic and Automation Control Conference (IMCEC)* (pp. 1519–1522).

Liu, W., Wang, Z., Liu, X., Zeng, N., Liu, Y., & Alsaadi, F. (2017). A survey of deep neural network architectures and their applications *Neurocomputing*, 234, 11–26.

Deng, L. (2014). A tutorial survey of architectures, algorithms, and applications for deep learning *APSIPA Transactions on Signal and Information Processing*, 3.

Grigorescu, S., Trasnea, B., Cocias, T., & Macesanu, G. (2019). A survey of deep learning techniques for autonomous driving *Journal of Field Robotics*.

De Mantaras, R., & Arcos, J. (2002). AI and music: From composition to expressive performance *AI magazine*, 23(3), 43–43.

Jiang, F., Jiang, Y., Zhi, H., Dong, Y., Li, H., Ma, S., Wang, Y., Dong, Q., Shen, H., & Wang, Y. (2017). Artificial intelligence in healthcare: past, present and future *Stroke and vascular neurology*, 2(4), 230–243.

Jain, A., Mao, J., & Mohiuddin, K. (1996). Artificial neural networks: A tutorial *Computer*, 29(3), 31–44.

Steven Abreu. (2019). Automated Architecture Design for Deep Neural Networks.

Lai Wei. (2018). Multi-hot Sparse Categorical Cross-entropy. <https://cwiki.apache.org/confluence/display/MXNET/Multi-hot+Sparse+Categorical+Cross-entropy>

Tensorflow. (2019). Tensor Board GitHub Repository. <https://github.com/tensorflow/tensorboard/blob/master/README.md>

Vitaly Bushaev. (2018). Adam — latest trends in deep learning optimization.
<https://towardsdatascience.com/adam-latest-trends-in-deep-learning-optimization-6be9a291375c>